

Large Language Models for Clinical Note Summarization

(2024)

Problem

Summarizing free-text clinical notes into structured discharge summaries.

Method

Fine-tuning a large language model on note-summary pairs.

Dataset

De-identified EHR clinical notes.

Evaluation

ROUGE score and physician review.

Results

Physician reviewers rated 82% of generated summaries as clinically acceptable without edits.

Limitations

This is an NLP-for-healthcare paper, not diagnostic imaging; included to test topic filtering against near-domain distractors.