

# **Benchmarking Classical Machine Learning vs Deep Learning for Tabular Clinical Risk Prediction**

(2022)

## **Problem**

Predicting 30-day hospital readmission risk from structured EHR data (not imaging).

## **Method**

Gradient boosted trees compared against a deep tabular model (TabNet).

## **Dataset**

MIMIC-IV structured clinical dataset.

## **Evaluation**

AUROC and calibration.

## **Results**

Gradient boosted trees slightly outperformed TabNet, 0.81 vs 0.79 AUROC.

## **Limitations**

Included here as a plausible but off-target match for an imaging-focused review; tests topic filtering.